

```

1  #pragma once
2  /*
3   * DBMS Implementation
4   *
5   * Contact: geopiskas@gmail.com
6   */
7
8
9  #include <stdio.h>
10 #include <stdlib.h>
11 #include <fcntl.h>
12 #include <string.h>
13 #include <vector>
14
15 #include "dbmsproj.h"
16 #include "recordOps.h"
17 #include "bufferOps.h"
18 #define GET_VARIABLE_NAME(Variable) (#Variable)
19
20 /*
21 * seed: seed to use in hash function
22 * buffer: buffer used, already loaded with a relation to hash
23 * size: the size in blocks of the relation loaded on buffer
24 * field: which field will be used for joining
25 *
26 * returns the pointer to the hash index
27 */
28 template <typename T> linkedRecordPtr** createHashIndex(char *seed, block_t<T> *buffer,
    unsigned int size, std::string const& field, int debugmode = 0) {
29     //initialize
30
31     // the hash index consists of a maximum of hashSize linked lists where
32     // each list has pointers to the records with common hash value
33
34     // the size of hashIndex
35     unsigned int hashSize = size * buffer->maxRecords;
36     linkedRecordPtr **hashIndex = (linkedRecordPtr**)malloc(hashSize * sizeof
    (linkedRecordPtr*));
37     for (unsigned int i = 0; i < hashSize; i++) {
38         hashIndex[i] = nullptr;
39     }
40
41     recordPtr start = newPtr(0,buffer->maxRecords);
42     unsigned int offset = (size - 1)*buffer->maxRecords + (buffer+size-1)->nreserved;
43     recordPtr end = newPtr(offset,buffer->maxRecords);
44
45     // starting from the very first record, all valid records in valid blocks
46     // are hashed
47     for (; start < end; incr(start,buffer->maxRecords)) {
48         if (!buffer[start.block].valid) {
49             start.record = buffer->maxRecords - 1;
50             continue;
51         }
52         T record = getRecord<T>(buffer, start);
53         if (record.valid) {
54             unsigned int index = hashRecord<T>(seed, record,hashSize,field);
55             linkedRecordPtr *ptr = (linkedRecordPtr*)malloc(sizeof(linkedRecordPtr));
56             ptr->ptr = start;

```

```

57     ptr->next = hashIndex[index];
58     hashIndex[index] = ptr;
59 }
60 }
61
62 // returns hashIndex
63 printf("Hash index is created successfully!\n");
64 return hashIndex;
65 };
66
67 /*
68 * infile: filename of the file whose records will be joined with the ones on buffer
69 * inBlocks: size of infile
70 * buffer: the buffer that is used (a file is already loaded on it)
71 * nmem_blocks: size of buffer
72 * size: the size of the file already loaded on buffer
73 * out: file descriptor of the outfile
74 * nres: number of pairs
75 * nios: number of ios
76 * field: which field will be used for joining
77 */
78
79 template <typename T1, typename T2> void hashAndProbe(std::vector<block_t<T1>> &r,
    block_t<T1> *buffer, block_t<T2>* buildbuffer,unsigned int size, std::string const&
    field, std::vector<block_t<T1>> &out, int debugmode = 0) {
80
81     unsigned int mod;
82
83     mod = size*buildbuffer->maxRecords;
84
85     char *seed = "agagagepiggeeq331516166fwhfsfrs";
86
87     // hash index for the records already on buffer is created
88     linkedRecordPtr **hashIndex = createHashIndex<T2>(seed, buildbuffer,size, field);
89     // pointer to the buffer block where blocks of infile are loaded
90     block_t<T1> *bufferIn = buffer;
91     // pointer to the last buffer block, where pairs for output are written
92     block_t<T1> *bufferOut = buffer + 1;
93
94
95     for (unsigned int i = 0; i < r.size(); i++) {
96         // if the block loaded is invalid, loads the next one
97         readBlocks<T1>(r, bufferIn, 1,i);
98         if (!(*bufferIn).valid) {
99             continue;
100         }
101         // each record of the loaded block is hashed
102         // then the linked list of the hash index for the corresponding hash value
103         // is examined, and if a record has same value as the current one, both
104         // are written to the output block
105
106         for (unsigned int j = 0; j < bufferIn->nreserved; j++) {
107             T1 record = (*bufferIn).entries[j];
108             if (record.valid) {
109                 unsigned int index = hashRecord<T1>(seed, record, mod,field);
110                 linkedRecordPtr *element = hashIndex[index];
111                 while (element) {
112                     T2 tmp = getRecord(buildbuffer, element->ptr);

```

```

113         if (compareRecords<T1,T2>(record, tmp, field) == 0) {
114             bufferOut->entries.push_back(record);
115             bufferOut->nreserved++;
116             //(*bufferOut).entries[(*bufferOut).nreserved++] = tmp;
117             //(*nres) += 1;
118
119             // if output block becomes full, writes it to the outfile
120             // and empties it
121
122             if (bufferOut->nreserved == bufferOut->maxRecords) {
123                 // writeBlocks(out, bufferOut, 1);
124                 out.push_back(*bufferOut);
125                 emptyBlock<T1>(bufferOut, bufferOut->maxRecords);
126                 bufferOut->blockid += 1;
127             }
128         }
129         element = element->next;
130     }
131 }
132 }
133 }
134 }
135
136 if (bufferOut->nreserved != 0) {
137     out.push_back(*bufferOut);
138 }
139 destroyHashIndex(hashIndex, size);
140 };
141 template <typename T1, typename T2> void hashAndProbefull(std::vector<block_t<T1>> &r,
142     block_t<T1> *buffer, block_t<T2> * buildbuffer, unsigned int size, std::string const&
143     field, std::vector<join_t<T1,T2>> &joinout, int debugmode = 0) {
144
145     printf("Building in-memory hash index and probing...\n");
146     unsigned int mod;
147     mod = size*buildbuffer->maxRecords;
148     char *seed = "agagagepiggeeq331516166fwhfsfrs";
149
150     // hash index for the records already on buffer is created
151     linkedRecordPtr **hashIndex = createHashIndex<T2>(seed, buildbuffer, size, field);
152     // pointer to the buffer block where blocks of infile are loaded
153     block_t<T1> *bufferIn = buffer;
154     // pointer to the last buffer block, where pairs for output are written
155     block_t<T1> *bufferOut = buffer + 1;
156
157     for (unsigned int i = 0; i < r.size(); i++) {
158         // if the block loaded is invalid, loads the next one
159         readBlocks<T1>(r, bufferIn, 1, i);
160         printf("Loading 1 block of proble relation into input buffer block.\n");
161         if (!(*bufferIn).valid) {
162             continue;
163         }
164         // each record of the loaded block is hashed
165         // then the linked list of the hash index for the corresponding hash value
166         // is examined, and if a record has same value as the current one, both
167         // are written to the output block
168
169         for (unsigned int j = 0; j < bufferIn->nreserved; j++) {
170             T1 record = (*bufferIn).entries[j];

```

```

170     if (record.valid) {
171         unsigned int index = hashRecord<T1>(seed, record, mod, field);
172         printf("Get 1 record from input buffer, hash it \n");
173         printf("Look up the hash index to match the records with same index=%d \n", index);
174         linkedRecordPtr *element = hashIndex[index];
175         while (element) {
176             T2 tmp = getRecord(buildbuffer, element->ptr);
177             if (compareRecords<T1, T2>(record, tmp, field) == 0) {
178                 bufferOut->entries.push_back(record);
179                 bufferOut->nreserved++;
180                 //(*bufferOut).entries[(*bufferOut).nreserved++] = tmp;
181                 //(*nres) += 1;
182                 // if output block becomes full, writes it to the outfile
183                 // and empties it
184                 printf("Join the two records. Get next record... \n");
185                 join_t<T1, T2> rec;
186                 rec.rec1 = record;
187                 rec.rec2 = tmp;
188                 joinout.push_back(rec);
189                 if (bufferOut->nreserved == bufferOut->maxRecords) {
190                     // writeBlocks(out, bufferOut, 1);
191                     printf("block is full, write back to disk. \n");
192                     emptyBlock<T1>(bufferOut, bufferOut->maxRecords);
193                     bufferOut->blockid += 1;
194                 }
195             }
196             element = element->next;
197         }
198     }
199 }
200 }
201 }
202 }
203 }
204 }
205
206 destroyHashIndex(hashIndex, size);
207 printf("End of probing. \n");
208 printf("===== \n");
209 printf("\n");
210 };
211 /*
212 * filename: the name of the file to be partitioned
213 * size: the size of the file
214 * seed: a seed for the hash function
215 * buffer: the buffer that is used
216 * nmem_blocks: size of buffer
217 * bucketFileNames: array with the filenames of the bucket files to be produced
218 * mod: to be used for hashing
219 * nios: number of ios
220 * field: which field will be used for joining
221 */
222 template<typename T> void createBuckets(std::vector<block_t<T>> &r, block_t<T> *buffer,
223     std::vector<std::vector<block_t<T>>> &partition, unsigned int mod, std::string const&
224     field, int debugmod = 0) {

```

```

225     unsigned int block_counts = r.size();
226     block_t<T> *bufferIn = buffer + MAX_MEMORY_BLOCKS - 1;
227
228     for (unsigned int i = 0; i < block_counts; i++)
229     {
230
231         *bufferIn = r[i]; //load one block from r into the last buffer block.
232         printf("Loading 1 block from disk into input buffer block. \n");
233         unsigned int max_records = bufferIn->entries.size();
234         for (unsigned int j = 0; j < max_records; j++)
235         {
236             T record = bufferIn->entries[j];
237             if (record.valid)
238             {
239                 unsigned int index = hashRecord<T>
240                     ("1235peqwtptquqewptuqptup1qtptu3421-58-12-35", record, mod, field);
241                 buffer[index].nreserved++;
242                 buffer[index].entries.push_back(record);
243                 //printf("record is hashed and put into %dth buffer block. \n", index);
244                 if (buffer[index].nreserved == buffer->maxRecords)
245                 {
246                     printf("%dth buffer block is full, write back to %dth partition. \n", index, index);
247                     buffer[index].valid = true;
248                     partition[index].push_back(buffer[index]);
249                     printf("Empty %dth buffer block. \n", index);
250                     emptyBlock<T>(buffer + index, buffer->maxRecords);
251                 }
252             }
253         }
254         printf("All the records in the input buffer have been processed, loading next block... \n");
255     }
256     printf("All the blocks in the disk have been processed.\n");
257     //put all non full blocks into the corresponding partition
258     for (unsigned int i = 0; i < mod; i++)
259     {
260         if (buffer[i].nreserved != 0)
261         {
262             buffer[i].valid = true;
263             partition[i].push_back(buffer[i]);
264             emptyBlock<T>(buffer + i, buffer[i].entries.size());
265         }
266     }
267
268
269     if (debugmod != 0)
270     {
271         printf("write all non-full buffer blocks back to corresponding partition. \n");
272         printf("The table has been partitioned into %d partitions: \n", mod);
273         printf("=====\n");
274         //display the results
275         unsigned int numPartition = 0;
276         unsigned int numBlocks=0;
277         unsigned int numRecords = 0;
278         for each (auto blocks in partition)

```

```

281     {
282     }
283     numBlocks = 0;
284     for each (auto block in blocks)
285     {
286         printf("The %dth Partition, %dth block:\n", numPartition, numBlocks);
287         block.printrecord();
288         numBlocks += 1;
289     }
290     numPartition += 1;
291     printf("  \n");
292 }
293 }
294 }
295 printf("End of partitioning the table. \n");
296 printf("===== \n");
297 printf("\n");
298 };
299
300
301 template <typename T1, typename T2> void SemiHashJoin(std::vector<block_t<T1>> &r,
std::vector<block_t<T2>> &s, char* field, std::vector<block_t<T1>> &out, int
debugmode=0) {
302
303     //figure out how many partitions to create
304     unsigned int smallSize;
305     if (r.size() > s.size()) smallSize = s.size(); else smallSize = r.size();
306     unsigned int mod=4;
307     /*mod = smallSize / (MAX_MEMORY_BLOCKS - 1);
308     if (smallSize % (MAX_MEMORY_BLOCKS - 1)) mod += 1;
309     if (mod > MAX_MEMORY_BLOCKS) mod = MAX_MEMORY_BLOCKS;
310     */
311     //create the partitions for r and s;
312     std::vector<std::vector<block_t<T1>>> r_partition(mod);
313     block_t<T1> T1buffer[MAX_MEMORY_BLOCKS];
314     std::vector<std::vector<block_t<T2>>> s_partition(mod);
315     block_t<T2> T2buffer[MAX_MEMORY_BLOCKS];
316     printf("Partitioning...\n");
317
318     createBuckets<T1>(r, T1buffer, r_partition, mod, field);
319     printf("The %s table has been partitioned into %d partitions: \n", GET_VARIABLE_NAME
(s), mod);
320     printf("===== \n");
321     createBuckets<T2>(s, T2buffer, s_partition, mod, field);
322
323     //process each partition each time
324     for (unsigned int i = 0; i < mod; i++)
325     {
326         unsigned int inblocks=r_partition[i].size();
327         block_t<T2> buildbuffer[MAX_MEMORY_BLOCKS-2];
328         block_t<T1> buffer[2];
329         if (s_partition[i].size()!=0)
330         {
331             //load one of s partitions into build buffer
332             for (unsigned int j = 0; j < s_partition[i].size(); j++)
333             {
334                 buildbuffer[j] = s_partition[i][j];
335             }
336             hashAndProbe(r_partition[i], buffer, buildbuffer, s_partition[i].size())

```

```

        field, out);
337     }
338 }
339 }
340 }
341
342 template <typename T1, typename T2> void HashJoin(std::vector<block_t<T1>> &r,
        std::vector<block_t<T2>> &s, char* field, std::vector<join_t<T1,T2>> &out, int
        debugmode = 0) {
343
344     //figure out how many partitions to create
345     unsigned int smallSize;
346     if (r.size() > s.size()) smallSize = s.size(); else smallSize = r.size();
347     unsigned int mod = 4;
348     /*mod = smallSize / (MAX_MEMORY_BLOCKS - 1);
349     if (smallSize % (MAX_MEMORY_BLOCKS - 1)) mod += 1;
350     if (mod > MAX_MEMORY_BLOCKS) mod = MAX_MEMORY_BLOCKS;
351     */
352     //create the partitions for r and s;
353     std::vector<std::vector<block_t<T1>>> r_partition(mod);
354     block_t<T1> T1buffer[MAX_MEMORY_BLOCKS];
355     std::vector<std::vector<block_t<T2>>> s_partition(mod);
356     block_t<T2> T2buffer[MAX_MEMORY_BLOCKS];
357     char* name = GET_VARIABLE_NAME(r);
358
359
360     printf("=====\n");
361
362     createBuckets<T1>(r, T1buffer, r_partition, mod, field);
363
364     printf("The %s table has been partitioned into %d partitions: \n", GET_VARIABLE_NAME
        (s), mod);
365     printf("=====\n");
366     createBuckets<T2>(s, T2buffer, s_partition, mod, field);
367
368     for (unsigned int i = 0; i < mod; i++)
369     {
370         unsigned int inblocks = r_partition[i].size();
371         block_t<T2> buildbuffer[MAX_MEMORY_BLOCKS - 2];
372         block_t<T1> buffer[2];
373
374         if (s_partition[i].size() != 0)
375         {
376             //load one of s partitions into build buffer
377             for (unsigned int j = 0; j < s_partition[i].size(); j++)
378             {
379                 buildbuffer[j] = s_partition[i][j];
380                 printf("Loading all blocks in %dth build partition into buffer. \n", j);
381             }
382             hashAndProbefull(r_partition[i], buffer, buildbuffer, s_partition[i].size(),
                field, out);
383         }
384     }
385 }
386 }

```