

```
1 // SimQP.cpp : Defines the entry point for the console application.
2 //
3
4 #include "stdafx.h"
5 #include "dbmsproj.h"
6 #include "semihashjoin.h"
7 #include <string>
8 #include <sstream>
9 #include <stdio.h>
10 #define GET_VARIABLE_NAME(Variable) (#Variable)
11 /**Using vectors represent disk to hold records block***/
12
13 //NYC - This site has a copy of branch, account and depositor tables.
14 std::vector<block_t<account_t>> account_NYC;
15 std::vector<block_t<branch_t>> branch_NYC;
16 std::vector<block_t<depositor_t>> depositor_NYC;
17
18 //SFO - This site has complete branch table and a fragment of account table.
19 //Account table will have only records that are local to SFO.
20 std::vector<block_t<branch_t>> branch_SFO;
21 std::vector<block_t<account_t>> frag_account_SFO;
22
23 //OMA - This site has fragments of Account and Depositor tables.
24 //These tables store records that are local to OMA.
25 std::vector<block_t<account_t>> frag_account_OMA;
26 std::vector<block_t<depositor_t>> frag_depositor_OMA;
27
28 //HOU (Houston) - This site has a complete copy of customer table.
29 std::vector<block_t<customer_t>> customer_HOU;
30
31
32 //unsigned int MAX_RECORDS_PER_BLOCK=7;
33
34
35 int main(int argc, char* argv[])
36 {
37     //initialize
38     char* filename;
39
40     //fill account table at NYC
41     filename = "Account_NYC.txt";
42     fillTable<account_t>(filename, account_NYC, 6);
43
44     //fill branch table at NYC
45     filename = "Branch_NYC.txt";
46     fillTable<branch_t>(filename, branch_NYC, 1);
47
48     //fill depositor table at NYC
49     filename = "Depositor_NYC.txt";
50     fillTable<depositor_t>(filename, depositor_NYC, 4);
51
52     //fill customer table at HOU
53     filename= "Customer_HOU.txt";
54     fillTable<customer_t>(filename, customer_HOU, 8);
55
56     //fill fragment of account table at SFO
57     filename = "Frag_Account_SFO.txt";
58     fillTable<account_t>(filename, frag_account_SFO, 2);
59 }
```

```

60
61 //fill fragment of depositor table at OMA
62 filename = "Frag_Depositor_OMA.txt";
63 fillTable<depositor_t>(filename, frag_depositor_OMA, 2);
64
65 //fill fragment of account table at OMA
66 filename = "Frag_Account_OMA.txt";
67 //fillTable<account_t>(filename, frag_account_OMA, 3);
68 //printf("All tables are filled.\n");
69 //printf("=====\n");
70 std::istringstream ss(argv[1]);
71 int query;
72 if (!(ss >> query)) std::cerr << "Invalid number" << argv[1] << std::endl;
73
74 std::vector<std::string> command;
75 std::vector<std::string> parameters;
76 std::string require_site;
77 std::string current_site;
78
79
80 switch (query)
81 {
82 case 1:
83     {
84         std::vector<block_t<customer_t>> out;
85         printf("Query 1: \n");
86         printf("Require site: NYC\n");
87         printf("Statement: Select* from customer |X depositor:\n");
88         printf("Processing:\n");
89         printf("The customer table is shipped from HOU to NYC site. \n");
90
91         printf("\n");
92         SemiHashJoin<customer_t,depositor_t>
93             (customer_HOU,depositor_NYC,"customer_name", out);
94         printf("The customer table are SemiHashJoined with depositor table at NYC
95             site. \n");
96         printf("Project all fields in customer table.\n");
97         printf("\n");
98         printf("Output: \n");
99         printf("=====\n");
100         printf("Name    Street  City \n");
101         printf("=====\n");
102         unsigned int count = 0;
103         for each ( auto block in out)
104         {
105             for each (auto record in block.entries)
106             {
107                 record.project(3, "customer_name",
108                     "customer_city","customer_street");
109                 count++;
110             }
111         }
112         printf("=====\n");
113         printf("The query output %d records.\n",count);
114         break;

```

```

115     case 2:
116     {
117
118         printf("Query 2: \n");
119         printf("Require site: SFO\n");
120         printf("Statement: Select name, balance from depositor |X| account where
121             branch_name='Chinatown':\n");
122         printf("Processing:\n");
123
124         printf("3) Ship the results from NYC to SFO and then do projection. \n");
125         printf("\n");
126         std::vector<block_t<account_t>> table_selec;
127         printf("Select the records with branch_name='Chinatown' using the fragment
128             of account table at SFO.\n");
129         select(account_NYC, table_selec, "branch_name", "Chinatown");
130
131         if (table_selec.size() == 0)
132         {
133             printf("no record is selected according to the condition.\n");
134             break;
135         }
136         else
137         {
138             unsigned int rec_selected = (table_selec.size() - 1)*table_selec
139                 [0].maxRecords + table_selec[table_selec.size()-1].nreserved;
140             //printf("%d records are selected.\n", rec_selected);
141
142             printf("Ship the selection results from SFO to HOU and hash joined with
143                 depositor table. \n");
144             std::vector<join_t<depositor_t,account_t>> joinout;
145             HashJoin<depositor_t,account_t>(depositor_NYC,
146                 table_selec,"account_number", joinout);
147             printf("The join results at HOU is shipped back to SFO. \n");
148             printf("\n");
149             printf("Output: \n");
150             printf("=====\n");
151             printf("Name Balance\n");
152             printf("=====\n");
153             unsigned int count=0;
154             for each (auto var in joinout)
155             {
156                 var.display(2,"customer_name","balance");
157                 count++;
158             }
159             printf("=====\n");
160             printf("The query output %d records.\n", count);
161             break;
162         }
163     }
164
165     case 3:
166     {
167
168         printf("Query 3: \n");
169         printf("Require site: SFO\n");
170         printf("Statement: Select street, city from customer|X (depositor |X account
171             where account_number='A10352'):\n");
172         printf("Processing:\n");
173
174         printf("\n");

```

```

167     std::vector<block_t<account_t>> table_selec;
168
169     printf("Search the fragment of account table at SFO to check if there have
170         records with account_number='A10352' \n");
171     select(frag_account_SF0, table_selec, "account_number", "A10352");
172     if (table_selec.size() == 0)
173     {
174         printf("no record is found according to the condition.\n");
175         select(account_NYC, table_selec, "account_number", "A10352");
176         printf("Search the account table at NYC.\n");
177     }
178     if (table_selec.size() == 0)
179     {
180         printf("no record is selected according to the condition.\n");
181         break;
182     }
183     else
184     {
185         unsigned int rec_selected = (table_selec.size() - 1)*table_selec
186             [0].maxRecords + table_selec[table_selec.size() - 1].nreserved;
187         //printf("%d records are selected.\n", rec_selected);
188     }
189
190     std::vector<block_t<depositor_t>> semi_out;
191     printf("The depositor table is semi-joined with intermediate result from select
192         operation at NYC.\n");
193     SemiHashJoin<depositor_t,account_t>(depositor_NYC,table_selec,"account_number",
194         semi_out);
195
196     printf("The intermediate results from semijoin are shipped from NYC to HOU. \n");
197     printf("The customer table at HOU is hash joined with the intermediate results.
198         \n");
199     std::vector<join_t<customer_t, depositor_t>> joinout;
200     HashJoin<customer_t, depositor_t>(customer_HOU, semi_out, "customer_name",
201         joinout);
202
203     printf("The final result is shipped from HOU to the SFO. \n");
204
205     printf("Output: \n");
206     printf("=====\n");
207     printf("Street      City\n");
208     printf("=====\n");
209
210     unsigned int count=0;
211     for each (auto var in joinout)
212     {
213         var.display(2, "customer_street", "customer_city");
214         count++;
215     }
216     printf("=====\n");
217     printf("The query output %d records.\n", count);
218     break;
219 }
220
221 case 4:
222 {
223
224     printf("Query 4: \n");

```

```

220     printf("Require site: NYC\n");
221     printf("Statement: Select account_number, balance, branch_name branch_city from
        branch|X| acocunt where account_number='A10352'):\n");
222     printf("Processing:\n");
223
224     std::vector<block_t<account_t>> table_selec;
225     printf("Select the record from the account table at NYC where
        account_number='A10352' \n");
226     select(account_NYC, table_selec, "account_number", "A10352");
227     if (table_selec.size() == 0)
228     {
229         //printf("no record is selected according to the condition.\n");
230         break;
231     }
232     else
233     {
234         unsigned int rec_selected = (table_selec.size() - 1)*table_selec
            [0].maxRecords + table_selec[table_selec.size() - 1].nreserved;
235         //printf("%d records are selected.\n", rec_selected);
236     }
237
238     std::vector<join_t<branch_t, account_t>> joinout;
239     printf("The branch table is hash joined with intermediate result from select
        operation at NYC.\n");
240     HashJoin<branch_t, account_t>(branch_NYC, table_selec, "branch_name", joinout);
241     printf("=====\n");
242
243     printf("Output: \n");
244     printf("=====\n");
245     printf("account_nmber    balance branch_name branch_city\n");
246     printf("=====\n");
247     unsigned int count = 0;
248     for each (auto var in joinout)
249     {
250         var.display(4, "account_number", "balance", "branch_name", "branch_city");
251         count++;
252     }
253     printf("=====\n");
254     printf("The query output %d records.\n", count);
255     break;
256 }
257 default:
258     break;
259 }
260
261 printf("\n");
262 }
263 }
264
265

```

```

1  #pragma once
2  /*
3   * DBMS Implementation
4   *
5   * Contact: geopiskas@gmail.com
6   */
7
8
9  #include <stdio.h>
10 #include <stdlib.h>
11 #include <fcntl.h>
12 #include <string.h>
13 #include <vector>
14
15 #include "dbmsproj.h"
16 #include "recordOps.h"
17 #include "bufferOps.h"
18 #define GET_VARIABLE_NAME(Variable) (#Variable)
19
20 /*
21 * seed: seed to use in hash function
22 * buffer: buffer used, already loaded with a relation to hash
23 * size: the size in blocks of the relation loaded on buffer
24 * field: which field will be used for joining
25 *
26 * returns the pointer to the hash index
27 */
28 template <typename T> linkedRecordPtr** createHashIndex(char *seed, block_t<T> *buffer,
    unsigned int size, std::string const& field, int debugmode = 0) {
29     //initialize
30
31     // the hash index consists of a maximum of hashSize linked lists where
32     // each list has pointers to the records with common hash value
33
34     // the size of hashIndex
35     unsigned int hashSize = size * buffer->maxRecords;
36     linkedRecordPtr **hashIndex = (linkedRecordPtr**)malloc(hashSize * sizeof
    (linkedRecordPtr*));
37     for (unsigned int i = 0; i < hashSize; i++) {
38         hashIndex[i] = nullptr;
39     }
40
41     recordPtr start = newPtr(0,buffer->maxRecords);
42     unsigned int offset = (size - 1)*buffer->maxRecords + (buffer+size-1)->nreserved;
43     recordPtr end = newPtr(offset,buffer->maxRecords);
44
45     // starting from the very first record, all valid records in valid blocks
46     // are hashed
47     for (; start < end; incr(start,buffer->maxRecords)) {
48         if (!buffer[start.block].valid) {
49             start.record = buffer->maxRecords - 1;
50             continue;
51         }
52         T record = getRecord<T>(buffer, start);
53         if (record.valid) {
54             unsigned int index = hashRecord<T>(seed, record,hashSize,field);
55             linkedRecordPtr *ptr = (linkedRecordPtr*)malloc(sizeof(linkedRecordPtr));
56             ptr->ptr = start;

```

```

57     ptr->next = hashIndex[index];
58     hashIndex[index] = ptr;
59 }
60 }
61
62 // returns hashIndex
63 printf("Hash index is created successfully!\n");
64 return hashIndex;
65 };
66
67 /*
68 * infile: filename of the file whose records will be joined with the ones on buffer
69 * inBlocks: size of infile
70 * buffer: the buffer that is used (a file is already loaded on it)
71 * nmem_blocks: size of buffer
72 * size: the size of the file already loaded on buffer
73 * out: file descriptor of the outfile
74 * nres: number of pairs
75 * nios: number of ios
76 * field: which field will be used for joining
77 */
78
79 template <typename T1, typename T2> void hashAndProbe(std::vector<block_t<T1>> &r,
    block_t<T1> *buffer, block_t<T2>* buildbuffer,unsigned int size, std::string const&
    field, std::vector<block_t<T1>> &out, int debugmode = 0) {
80
81     unsigned int mod;
82
83     mod = size*buildbuffer->maxRecords;
84
85     char *seed = "agagagepiggeeq331516166fwhfsfrs";
86
87     // hash index for the records already on buffer is created
88     linkedRecordPtr **hashIndex = createHashIndex<T2>(seed, buildbuffer,size, field);
89     // pointer to the buffer block where blocks of infile are loaded
90     block_t<T1> *bufferIn = buffer;
91     // pointer to the last buffer block, where pairs for output are written
92     block_t<T1> *bufferOut = buffer + 1;
93
94
95     for (unsigned int i = 0; i < r.size(); i++) {
96         // if the block loaded is invalid, loads the next one
97         readBlocks<T1>(r, bufferIn, 1,i);
98         if (!(*bufferIn).valid) {
99             continue;
100         }
101         // each record of the loaded block is hashed
102         // then the linked list of the hash index for the corresponding hash value
103         // is examined, and if a record has same value as the current one, both
104         // are written to the output block
105
106         for (unsigned int j = 0; j < bufferIn->nreserved; j++) {
107             T1 record = (*bufferIn).entries[j];
108             if (record.valid) {
109                 unsigned int index = hashRecord<T1>(seed, record, mod,field);
110                 linkedRecordPtr *element = hashIndex[index];
111                 while (element) {
112                     T2 tmp = getRecord(buildbuffer, element->ptr);

```

```

113         if (compareRecords<T1,T2>(record, tmp, field) == 0) {
114             bufferOut->entries.push_back(record);
115             bufferOut->nreserved++;
116             //(*bufferOut).entries[(*bufferOut).nreserved++] = tmp;
117             //(*nres) += 1;
118
119             // if output block becomes full, writes it to the outfile
120             // and empties it
121
122             if (bufferOut->nreserved == bufferOut->maxRecords) {
123                 // writeBlocks(out, bufferOut, 1);
124                 out.push_back(*bufferOut);
125                 emptyBlock<T1>(bufferOut, bufferOut->maxRecords);
126                 bufferOut->blockid += 1;
127             }
128         }
129         element = element->next;
130     }
131 }
132 }
133 }
134 }
135
136 if (bufferOut->nreserved != 0) {
137     out.push_back(*bufferOut);
138 }
139 destroyHashIndex(hashIndex, size);
140 };
141 template <typename T1, typename T2> void hashAndProbefull(std::vector<block_t<T1>> &r,
142     block_t<T1> *buffer, block_t<T2> * buildbuffer, unsigned int size, std::string const&
143     field, std::vector<join_t<T1,T2>> &joinout, int debugmode = 0) {
144
145     printf("Building in-memory hash index and probing...\n");
146     unsigned int mod;
147     mod = size*buildbuffer->maxRecords;
148     char *seed = "agagagepiggeeq331516166fwhfsfrs";
149
150     // hash index for the records already on buffer is created
151     linkedRecordPtr **hashIndex = createHashIndex<T2>(seed, buildbuffer, size, field);
152     // pointer to the buffer block where blocks of infile are loaded
153     block_t<T1> *bufferIn = buffer;
154     // pointer to the last buffer block, where pairs for output are written
155     block_t<T1> *bufferOut = buffer + 1;
156
157     for (unsigned int i = 0; i < r.size(); i++) {
158         // if the block loaded is invalid, loads the next one
159         readBlocks<T1>(r, bufferIn, 1, i);
160         printf("Loading 1 block of proble relation into input buffer block.\n");
161         if (!(*bufferIn).valid) {
162             continue;
163         }
164         // each record of the loaded block is hashed
165         // then the linked list of the hash index for the corresponding hash value
166         // is examined, and if a record has same value as the current one, both
167         // are written to the output block
168
169         for (unsigned int j = 0; j < bufferIn->nreserved; j++) {
170             T1 record = (*bufferIn).entries[j];

```



```

170         if (record.valid) {
171             unsigned int index = hashRecord<T1>(seed, record, mod, field);
172             printf("Get 1 record from input buffer, hash it \n");
173             printf("Look up the hash index to match the records with same index=%d \n", index);
174             linkedRecordPtr *element = hashIndex[index];
175             while (element) {
176                 T2 tmp = getRecord(buildbuffer, element->ptr);
177                 if (compareRecords<T1, T2>(record, tmp, field) == 0) {
178                     bufferOut->entries.push_back(record);
179                     bufferOut->nreserved++;
180                     //(*bufferOut).entries[(*bufferOut).nreserved++] = tmp;
181                     //(*nres) += 1;
182                     // if output block becomes full, writes it to the outfile
183                     // and empties it
184                     printf("Join the two records. Get next record... \n");
185                     join_t<T1, T2> rec;
186                     rec.rec1 = record;
187                     rec.rec2 = tmp;
188                     joinout.push_back(rec);
189                     if (bufferOut->nreserved == bufferOut->maxRecords) {
190                         // writeBlocks(out, bufferOut, 1);
191                         printf("block is full, write back to disk. \n");
192                         emptyBlock<T1>(bufferOut, bufferOut->maxRecords);
193                         bufferOut->blockid += 1;
194                     }
195                     element = element->next;
196                 }
197             }
198         }
199     }
200 }
201
202
203
204
205
206 destroyHashIndex(hashIndex, size);
207 printf("End of probing. \n");
208 printf("===== \n");
209 printf("\n");
210 };
211 /*
212 * filename: the name of the file to be partitioned
213 * size: the size of the file
214 * seed: a seed for the hash function
215 * buffer: the buffer that is used
216 * nmem_blocks: size of buffer
217 * bucketFileNames: array with the filenames of the bucket files to be produced
218 * mod: to be used for hashing
219 * nios: number of ios
220 * field: which field will be used for joining
221 */
222 template<typename T> void createBuckets(std::vector<block_t<T>> &r, block_t<T> *buffer,
223     std::vector<std::vector<block_t<T>>> &partition, unsigned int mod, std::string const&
224     field, int debugmod = 0) {

```

```

225     unsigned int block_counts = r.size();
226     block_t<T> *bufferIn = buffer + MAX_MEMORY_BLOCKS - 1;
227
228     for (unsigned int i = 0; i < block_counts; i++)
229     {
230
231         *bufferIn = r[i]; //load one block from r into the last buffer block.
232         printf("Loading 1 block from disk into input buffer block. \n");
233         unsigned int max_records = bufferIn->entries.size();
234         for (unsigned int j = 0; j < max_records; j++)
235         {
236             T record = bufferIn->entries[j];
237             if (record.valid)
238             {
239                 unsigned int index = hashRecord<T>
240                     ("1235peqwtptquqewptuqptup1qtptu3421-58-12-35", record, mod, field);
241                 buffer[index].nreserved++;
242                 buffer[index].entries.push_back(record);
243                 //printf("record is hashed and put into %dth buffer block. \n", index);
244                 if (buffer[index].nreserved == buffer->maxRecords)
245                 {
246                     printf("%dth buffer block is full, write back to %dth partition. \n", index, index);
247                     buffer[index].valid = true;
248                     partition[index].push_back(buffer[index]);
249                     printf("Empty %dth buffer block. \n", index);
250                     emptyBlock<T>(buffer + index, buffer->maxRecords);
251                 }
252             }
253         }
254         printf("All the records in the input buffer have been processed, loading next block... \n");
255     }
256     printf("All the blocks in the disk have been processed.\n");
257     //put all non full blocks into the corresponding partition
258     for (unsigned int i = 0; i < mod; i++)
259     {
260         if (buffer[i].nreserved != 0)
261         {
262             buffer[i].valid = true;
263             partition[i].push_back(buffer[i]);
264             emptyBlock<T>(buffer + i, buffer[i].entries.size());
265         }
266     }
267
268     if (debugmod != 0)
269     {
270         printf("write all non-full buffer blocks back to corresponding partition. \n");
271         printf("The table has been partitioned into %d partitions: \n", mod);
272         printf("=====\n");
273         //display the results
274         unsigned int numPartition = 0;
275         unsigned int numBlocks=0;
276         unsigned int numRecords = 0;
277         for each (auto blocks in partition)

```

```

281     {
282     }
283     numBlocks = 0;
284     for each (auto block in blocks)
285     {
286         printf("The %dth Partition, %dth block:\n", numPartition, numBlocks);
287         block.printrecord();
288         numBlocks += 1;
289     }
290     numPartition += 1;
291     printf("  \n");
292 }
293 }
294 }
295 printf("End of partitioning the table. \n");
296 printf("===== \n");
297 printf("\n");
298 };
299
300
301 template <typename T1, typename T2> void SemiHashJoin(std::vector<block_t<T1>> &r,
std::vector<block_t<T2>> &s, char* field, std::vector<block_t<T1>> &out, int
debugmode=0) {
302
303     //figure out how many partitions to create
304     unsigned int smallSize;
305     if (r.size() > s.size()) smallSize = s.size(); else smallSize = r.size();
306     unsigned int mod=4;
307     /*mod = smallSize / (MAX_MEMORY_BLOCKS - 1);
308     if (smallSize % (MAX_MEMORY_BLOCKS - 1)) mod += 1;
309     if (mod > MAX_MEMORY_BLOCKS) mod = MAX_MEMORY_BLOCKS;
310     */
311     //create the partitions for r and s;
312     std::vector<std::vector<block_t<T1>>> r_partition(mod);
313     block_t<T1> T1buffer[MAX_MEMORY_BLOCKS];
314     std::vector<std::vector<block_t<T2>>> s_partition(mod);
315     block_t<T2> T2buffer[MAX_MEMORY_BLOCKS];
316     printf("Partitioning...\n");
317
318     createBuckets<T1>(r, T1buffer, r_partition, mod, field);
319     printf("The %s table has been partitioned into %d partitions: \n", GET_VARIABLE_NAME
(s), mod);
320     printf("===== \n");
321     createBuckets<T2>(s, T2buffer, s_partition, mod, field);
322
323     //process each partition each time
324     for (unsigned int i = 0; i < mod; i++)
325     {
326         unsigned int inblocks=r_partition[i].size();
327         block_t<T2> buildbuffer[MAX_MEMORY_BLOCKS-2];
328         block_t<T1> buffer[2];
329         if (s_partition[i].size()!=0)
330         {
331             //load one of s partitions into build buffer
332             for (unsigned int j = 0; j < s_partition[i].size(); j++)
333             {
334                 buildbuffer[j] = s_partition[i][j];
335             }
336             hashAndProbe(r_partition[i], buffer, buildbuffer, s_partition[i].size())

```

```

        field, out);
337     }
338 }
339 }
340 }
341
342 template <typename T1, typename T2> void HashJoin(std::vector<block_t<T1>> &r,
        std::vector<block_t<T2>> &s, char* field, std::vector<join_t<T1,T2>> &out, int
        debugmode = 0) {
343
344     //figure out how many partitions to create
345     unsigned int smallSize;
346     if (r.size() > s.size()) smallSize = s.size(); else smallSize = r.size();
347     unsigned int mod = 4;
348     /*mod = smallSize / (MAX_MEMORY_BLOCKS - 1);
349     if (smallSize % (MAX_MEMORY_BLOCKS - 1)) mod += 1;
350     if (mod > MAX_MEMORY_BLOCKS) mod = MAX_MEMORY_BLOCKS;
351     */
352     //create the partitions for r and s;
353     std::vector<std::vector<block_t<T1>>> r_partition(mod);
354     block_t<T1> T1buffer[MAX_MEMORY_BLOCKS];
355     std::vector<std::vector<block_t<T2>>> s_partition(mod);
356     block_t<T2> T2buffer[MAX_MEMORY_BLOCKS];
357     char* name = GET_VARIABLE_NAME(r);
358
359
360     printf("=====\n");
361
362     createBuckets<T1>(r, T1buffer, r_partition, mod, field);
363
364     printf("The %s table has been partitioned into %d partitions: \n", GET_VARIABLE_NAME
        (s), mod);
365     printf("=====\n");
366     createBuckets<T2>(s, T2buffer, s_partition, mod, field);
367
368     for (unsigned int i = 0; i < mod; i++)
369     {
370         unsigned int inblocks = r_partition[i].size();
371         block_t<T2> buildbuffer[MAX_MEMORY_BLOCKS - 2];
372         block_t<T1> buffer[2];
373
374         if (s_partition[i].size() != 0)
375         {
376             //load one of s partitions into build buffer
377             for (unsigned int j = 0; j < s_partition[i].size(); j++)
378             {
379                 buildbuffer[j] = s_partition[i][j];
380                 printf("Loading all blocks in %dth build partition into buffer. \n", j);
381             }
382             hashAndProbefull(r_partition[i], buffer, buildbuffer, s_partition[i].size(),
                field, out);
383         }
384     }
385 }
386 }

```

```

1  /*
2  * DBMS Implementation
3  */
4
5  #ifndef _DBMSPROJ_H
6  #define _DBMSPROJ_H
7
8  #include "stdafx.h"
9  #include <fstream>
10 #include <string>
11 #include <stdio.h>
12 #include <iostream>
13 #include <sstream>
14 #include <vector>
15 #include <cstdint>
16 // #include "bufferOps.h"
17
18 #define TUPLES_PER_ACCOUNT_BLOCK 10
19 #define TUPLES_PER_BRANCH_BLOCK 7
20 #define TUPLES_PER_CUSTOMER_BLOCK 8
21 #define TUPLES_PER_DEPOSITOR_BLOCK 15
22
23 #define NUM_BLOCKS_ACCOUNT 6
24 #define NUM_BLOCKS_BRANCH 1
25 #define NUM_BLOCKS_CUSTOMER 8
26 #define NUM_BLOCKS_DEPOSITOR 4
27
28 #define MAX_RECORDS_PER_BUCKET 20
29
30
31 #define MAX_MEMORY_BLOCKS 5
32
33 //definition of record
34 //structure of records in account table;
35 enum field_code {e_account_number,e_branch_name,e_balance,e_branch_city,e_assets,e_customer_name,e_customer_street,e_customer_city};
36 inline field_code getfield(std::string const& field)
37 {
38     if (field == "account_number") return e_account_number;
39     if (field == "branch_name") return e_branch_name;
40     if (field == "balance") return e_balance;
41     if (field == "branch_city") return e_branch_city;
42     if (field == "assets") return e_assets;
43     if (field == "customer_name") return e_customer_name;
44     if (field == "customer_street") return e_customer_street;
45     if (field == "customer_city") return e_customer_city;
46 };
47 typedef struct {
48     std::string account_number;
49     std::string branch_name;
50     int balance = 0;
51     bool valid=false; //if set, then the record is valid
52     void setCol(std::string &val1, std::string &val2, std::string &val3)
53     {
54         account_number = val1;
55         branch_name = val2;
56         balance = std::stoi(val3,0);

```

```

57     valid = true;
58 }
59 std::string getCol(std::string const& field)
60 {
61     field_code fields = getfield(field);
62
63     switch (fields)
64     {
65     case e_account_number: return account_number; break;
66     case e_branch_name: return branch_name; break;
67     case e_balance: return std::to_string(balance); break;
68     default: return (" "); break;
69     }
70 }
71
72 void display() {
73     std::cout << account_number << " " << branch_name << " " << balance <<
74         std::endl;
75 }
76 void project(int num, const char* field, ...)
77 {
78     std::string fields(field);
79     if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
80     va_list arguments;
81     va_start(arguments, field);
82     for (size_t i = 0; i < num - 1; i++)
83     {
84         std::string fields(va_arg(arguments, char*));
85         if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
86     }
87     std::cout << std::endl;
88     va_end(arguments);
89 }
90 }account_t;
91
92 //structure of records in branch table;
93 typedef struct {
94     std::string branch_name;
95     std::string branch_city;
96     int assets = 0;
97     bool valid=false; //if set, then the record is valid
98     std::string getCol(std::string const& field)
99     {
100         field_code fields = getfield(field);
101
102         switch (fields)
103         {
104         case e_branch_name: return branch_name; break;
105         case e_branch_city: return branch_city; break;
106         //case 3: return assets; break;
107         default: return (" ");
108             break;
109         }
110     }
111 void setCol(std::string &val1, std::string &val2, std::string &val3)
112 {
113     branch_name = val1;
114     branch_city = val2;

```

```

115     assets = std::stoi(val3, 0);
116     valid = true;
117 }
118 void display() {
119     std::cout << branch_name << " " << branch_city << " " << assets << std::endl;
120 }
121 void project(int num, const char* field, ...)
122 {
123     std::string fields(field);
124     if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
125     va_list arguments;
126     va_start(arguments, field);
127     for (size_t i = 0; i < num - 1; i++)
128     {
129         std::string fields(va_arg(arguments, char*));
130         if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
131     }
132     std::cout << std::endl;
133     va_end(arguments);
134 }
135 }branch_t;
136
137 //structure of records in customer table;
138 typedef struct {
139     std::string customer_name;
140     std::string customer_street;
141     std::string customer_city;
142     bool valid=false; //if set, then the record is valid
143     void setCol(std::string &val1, std::string &val2, std::string &val3)
144     {
145         customer_name = val1;
146         customer_street = val2;
147         customer_city = val3;
148         valid = true;
149     }
150     std::string getCol(std::string const& field)
151     {
152         field_code fields = getfield(field);
153         switch (fields)
154         {
155             case e_customer_name: return customer_name; break;
156             case e_customer_street: return customer_street; break;
157             case e_customer_city: return customer_city; break;
158             default: return (" ");
159             break;
160         }
161     }
162     void display() {
163         std::cout << customer_name << " " << customer_street << " " << customer_city << "
164         std::endl;
165     }
166     void project(int num, const char* field, ...)
167     {
168         std::string fields(field);
169         if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
170         va_list arguments;
171         va_start(arguments, field);
172         for (size_t i = 0; i < num - 1; i++)

```

```

172     {
173         std::string fields(va_arg(arguments, char*));
174         if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
175     }
176     std::cout << std::endl;
177     va_end(arguments);
178 }
179 }customer_t;
180
181 //structure of records in depositor table;
182 typedef struct {
183     std::string customer_name;
184     std::string account_number;
185     bool valid=false; //if set, then the record is valid
186     void setCol(std::string &val1, std::string &val2, std::string &val3)
187     {
188         customer_name = val1;
189         account_number = val2;
190         valid = true;
191     }
192
193     std::string getCol(std::string const& field)
194     {
195         field_code fields = getfield(field);
196         switch (fields)
197         {
198             case e_customer_name: return customer_name; break;
199             case e_account_number: return account_number; break;
200             default: return (" "); break;
201         }
202     }
203     void display() {
204         std::cout << account_number << " " << customer_name<< std::endl;
205     }
206     void project(int num, const char* field, ...)
207     {
208         std::string fields(field);
209         if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
210         va_list arguments;
211         va_start(arguments, field);
212         for (size_t i = 0; i < num - 1; i++)
213         {
214             std::string fields(va_arg(arguments, char*));
215             if (getCol(fields) != (" ")) std::cout << getCol(fields) << " ";
216         }
217         std::cout << std::endl;
218         va_end(arguments);
219     }
220 }depositor_t;
221
222 template<typename T> struct block_t {
223     unsigned int blockid;
224     unsigned int nreserved;
225     unsigned int maxRecords;
226     bool valid;
227     std::vector<T> entries;
228
229     block_t()
230     {

```



```

231     blockid = 0;
232     nreserved = 0;
233     valid = false;
234     if (std::is_same<T, account_t>::value) maxRecords = 10;
235     if (std::is_same<T, branch_t>::value) maxRecords = 7;
236     if (std::is_same<T, customer_t>::value) maxRecords = 8;
237     if (std::is_same<T, depositor_t>::value) maxRecords = 15;
238     //entries.reserve(maxRecords);
239 }
240
241 void printrecord()
242 {
243     for (unsigned int i = 0; i < entries.size(); i++)
244     {
245         entries[i].display();
246     }
247 }
248
249 };
250
251
252
253
254 template<typename T1, typename T2> struct join_t {
255     T1 rec1;
256     T2 rec2;
257     void display(int num, std::string field, ...)
258     {
259         std::string fields(field);
260         if (rec1.getCol(fields) != (" ")) std::cout << rec1.getCol(fields) << " ";
261         if (rec2.getCol(fields) != (" ")) std::cout << rec2.getCol(fields) << " ";
262         va_list arguments;
263         va_start (arguments, field);
264         for (size_t i = 0; i < num-1; i++)
265         {
266             std::string fields(va_arg(arguments, char*));
267             if (rec1.getCol(fields) != (" ")) std::cout << rec1.getCol(fields) << " ";
268             if (rec2.getCol(fields) != (" ")) std::cout << rec2.getCol(fields) << " ";
269         }
270         std::cout << std::endl;
271         va_end(arguments);
272     }
273 };
274
275
276
277 /*
278
279
280
281
282 template<typename T> void fillTable(char* filename, std::vector<block_t<T>> &table,
283     unsigned int numblocks, int debugmode = 0) {
284
285     std::ifstream file(filename);

```

```

286     int rec_counts = 0;
287     int block_counts = 0;
288
289     for (unsigned int i = 0; i < numblocks; i++)
290     {
291         std::string line;
292         unsigned int index = 0;
293         block_t<T> block;
294
295         while ((index < block.maxRecords) && (std::getline(file, line)))
296         {
297             std::string val1;
298             std::string val2;
299             std::string val3;
300             T rec;
301             std::stringstream linestream(line);
302             std::string data;
303             std::getline(linestream, val1, '\t');
304             std::getline(linestream, val2, '\t');
305             std::getline(linestream, val3, '\t');
306
307             rec.setCol(val1, val2, val3);
308             rec.valid = true;
309             block.entries.push_back(rec);
310             rec_counts++;
311             index++;
312         }
313         block.blockid = i;
314         block.nreserved = index;
315         block.valid = true;
316         table.push_back(block);
317         block_counts++;
318     }
319
320     if (debugmode != 0)
321     {
322
323         printf("%s Table in NYC site has %d records in %d blocks.\n", filename, rec_counts,
324             block_counts);
325         for each (block_t<T> block in table)
326         {
327             block.printrecord();
328             printf("\n");
329         }
330         printf("\n");
331     }
332     template <typename T> void select(std::vector<block_t<T>> &in, std::vector<block_t<T>>
333         &out, std::string const & field, std::string const & val, int debugmod=0)
334     {
335         printf("Start processing selection.... \n");
336         block_t<T> buffer[MAX_MEMORY_BLOCKS];
337         //figure out how many times need to load the blocks for the input relation
338         unsigned int size = ((in.size() + MAX_MEMORY_BLOCKS - 2)) / (MAX_MEMORY_BLOCKS - 1);
339
340         //set the out block pointer to last block of the buffer
341         block_t<T> *bufferOut = buffer + MAX_MEMORY_BLOCKS - 1;
342
343         //process the blocks loaded into the buffer each time

```

```
343     for (unsigned int i = 0; i < size; i++)
344     {
345         int recordcounts = 0;
346         int nblocks=readBlocks<T>(in, buffer, (MAX_MEMORY_BLOCKS - 1), i*
347             (MAX_MEMORY_BLOCKS - 1));
348         recordPtr start = newPtr(0, buffer->maxRecords);
349         //calculate the total records of the loaded blocks in the buffer
350         unsigned int offset = (nblocks - 1)*buffer->maxRecords + (buffer + nblocks - 1)-
351             >nreserved;
352         recordPtr end = newPtr(offset, buffer->maxRecords);
353         // starting from the very first record, all valid records in valid blocks are
354         hashed
355         for (; start < end; incr(start, buffer->maxRecords))
356         {
357             if (!buffer[start.block].valid) {
358                 start.record = buffer->maxRecords - 1;
359                 continue;
360             }
361             T record = getRecord<T>(buffer, start);
362             if (record.getCol(field) == val)
363             {
364                 recordcounts++;
365                 bufferOut->entries.push_back(record);
366                 bufferOut->nreserved++;
367                 if (bufferOut->nreserved == bufferOut->maxRecords)
368                 {
369                     out.push_back(*bufferOut);
370                     emptyBlock<T>(bufferOut, bufferOut->maxRecords);
371                 }
372             }
373         }
374         printf("%d records are selected... \n",recordcounts);
375         //if there has any records left in the output buffer, push into the out vector
376         if (bufferOut->nreserved != 0) out.push_back(*bufferOut);
377         printf("End of selection operation.... \n");
378     }
379 }
380
381 #endif
382
```

```

1  /*
2  * DBMS Implementation
3  */
4
5  #ifndef BUFFEROPS_H
6  #define BUFFEROPS_H
7
8  #include <io.h>
9  #include <fcntl.h>
10 #include <sys/types.h>
11
12 #include "dbmsproj.h"
13 #include<iostream>
14
15
16 // empties a block
17 template<typename T> void emptyBlock(block_t<T> *buffer, unsigned int size) {
18     buffer->nreserved = 0;
19     buffer->entries.clear();
20     //printf("empty buffer block... \n");
21 };
22
23
24 // empties the whole buffer
25
26 template<typename T> void emptyBuffer(block_t<T> *buffer, unsigned int size) {
27     for (uint i = 0; i < size; i++) {
28         emptyBlock(buffer + i);
29         buffer[i].valid = true;
30     }
31     printf("empty buffer ...");
32 };
33
34 // opens filename for writing (append mode), and writes size blocks
35 // starting from pointer buffer
36
37 template<typename T> unsigned int writeBlocks(block_t<T> &relation, block_t<T> &buffer,  ➤
    unsigned int offset, unsigned int size) {
38     unsigned int writecounts = 0;
39
40     for (size_t i = 0; i < size; i++)
41     {
42         relation[i+offset] = buffer[i];
43         writecounts++;
44     }
45     printf("Write %d blocks into disk...\n", readcount);
46     return size;
47 };
48
49 // reads size blocks to buffer
50
51 template <typename T> unsigned int readBlocks(std::vector<block_t<T>>& relation,  ➤
    block_t<T>* buffer, unsigned int size) {
52     unsigned int readcount = 0;
53     for (unsigned int i = 0; i < size; i++)
54     {
55         if (i<relation.size())
56         {
57             buffer[i] = relation[offset + i];
58             readcount++;

```

```
59     }
60     else
61     {
62         break;
63     }
64 }
65 printf("loading %d blocks into buffer...\n",readcount);
66 return readcount;
67 };
68 template <typename T> unsigned int readBlocks(std::vector<block_t<T>>& relation,
        block_t<T>* buffer,unsigned int size,unsigned int offset) {
69     unsigned int readcount = 0;
70     for (unsigned int i = 0; i < size; i++)
71     {
72         if ((offset+i)<relation.size())
73         {
74             buffer[i] = relation[offset+i];
75             readcount++;
76         }
77         else
78         {
79             break;
80         }
81     }
82     printf("loading %d blocks into buffer...\n",readcount);
83     return readcount;
84 };
85 #endif
86
87
```

```
1  /*
2  * DBMS Implementation
3  */
4
5  #ifndef RECORDPTR_H
6  #define RECORDPTR_H
7
8  #include <sys/types.h>
9  #include "dbmsproj.h"
10
11 // location of a specific record in the buffer
12 // block is the index number of the block the record is in
13 // record is the index number of the record in the entries table
14
15 typedef struct {
16     unsigned int block;
17     unsigned int record;
18 } recordPtr;
19
20 // struct for making linked lists of recordPtrs
21
22 struct linkedRecordPtr {
23     recordPtr ptr;
24     linkedRecordPtr *next;
25 };
26
27 //overloading operators for use with struct recordPtr
28 inline bool operator==(const recordPtr &ptr1, const recordPtr &ptr2) {
29     if (ptr1.block==ptr2.block && ptr1.record==ptr2.record){
30         return true;
31     }
32     else return false;
33 };
34 inline bool operator<(const recordPtr &ptr1, const recordPtr &ptr2) {
35     if (ptr1.block == ptr2.block) {
36         if (ptr1.record < ptr2.record)
37         {
38             return true;
39         }
40         else
41         {
42             return false;
43         }
44     }
45     else
46     {
47         if (ptr1.block < ptr2.block)
48         {
49             return true;
50         }
51         else
52         {
53             return false;
54         }
55     }
56 }
57
58 inline bool operator<=(const recordPtr &ptr1, const recordPtr &ptr2)
59 {
60     if (ptr1 < ptr2 || ptr1 == ptr2) { return true; }
```

```
61     else return false;
62 }
63
64 inline recordPtr newPtr(recordPtr ptr,unsigned int offset,unsigned int size) {
65     recordPtr result;
66     result.block = ptr.block + offset / size;
67     int rest = offset % size;
68
69     if (ptr.record + rest >= size) {
70         result.block += 1;
71         result.record = ptr.record + rest - size;
72     }
73     else {
74         result.record = ptr.record + rest;
75     }
76     return result;
77 }
78
79 inline recordPtr newPtr(unsigned int offset,unsigned int size) {
80     recordPtr zero;
81     zero.block = 0;
82     zero.record = 0;
83     return newPtr(zero,offset,size);
84 }
85
86 // increases the ptr so that it points to the next record
87 // if pointing at the end of a block, moves to the start of the next
88
89 inline void incr(recordPtr &ptr,unsigned int size) {
90     if (ptr.record < size - 1) {
91         ptr.record += 1;
92     } else {
93         ptr.record = 0;
94         ptr.block += 1;
95     }
96 }
97
98 // decreases the ptr so that it points to the next record
99 // if pointing at the start of a block, moves to the end of the previous
100
101 inline void decr(recordPtr &ptr,unsigned int size) {
102     if (ptr.record > 0) {
103         ptr.record -= 1;
104     } else if (ptr.block > 0) {
105         ptr.record = size - 1;
106         ptr.block -= 1;
107     }
108 }
109
110 #endif
```

```

1  /*
2  * DBMS Implementation
3  */
4
5  #ifndef RECORDOPS_H
6  #define RECORDOPS_H
7
8  #include <string.h>
9
10 #include "dbmsproj.h"
11 #include "recordPtr.h"
12
13 // given a buffer and a recordPtr, returns corresponding record
14
15 template <typename T> T getRecord(block_t<T> *buffer, recordPtr ptr) {
16     return buffer[ptr.block].entries[ptr.record];
17 };
18
19 // given a buffer, a record and a recordPtr, places the record where recordPtr points
20
21 template <typename T> void setRecord(block_t<T> *buffer, T rec, recordPtr ptr) {
22     buffer[ptr.block].entries[ptr.record] = rec;
23 };
24
25 // given a buffer and 2 recordPtrs, swaps the records where ptrs point
26
27 template <typename T> void swapRecords(block_t<T> *buffer, recordPtr ptr1, recordPtr ptr2) {
28     {
29         T tmp = getRecord(buffer, ptr1);
30         setRecord(buffer, getRecord(buffer, ptr2), ptr1);
31         setRecord(buffer, tmp, ptr2);
32     }
33 };
34
35 // given 2 records and field, compares them
36 // -1 is returned if rec1 has lower field value than rec2
37 // 0 is returned if rec1 and rec2 have equal field values
38 // 1 is returned if rec1 has higher field value than rec2
39 template <typename T1, typename T2> int compareRecords(T1 &rec1, T2 &rec2, std::string
40     const& field) {
41     if (rec1.getCol(field) == rec2.getCol(field)) { /*printf("Two records matched! \n");*/
42         return 0; }
43     else return -1;
44 }
45
46 // hash function for integers
47
48 inline unsigned int hashInt(unsigned int num, unsigned int mod, unsigned int seed) {
49     num += seed;
50     num = (num + 0x7ed55d16) + (num << 12);
51     num = (num^0xc761c23c) ^ (num >> 19);
52     num = (num + 0x165667b1) + (num << 5);
53     num = (num + 0xd3a2646c) ^ (num << 9);
54     num = (num + 0xfd7046c5) + (num << 3);
55     num = (num^0xb55a4f09) ^ (num >> 16);
56     return num % mod;
57 }
58
59 // hash function for strings
60

```



```
57 inline unsigned int hashString(std::string str_text, unsigned int mod, unsigned int seed)
    {
58
59     char *str = new char[str_text.length() + 1];
60
61     std::strcpy(str, str_text.c_str());
62     str[str_text.length()] = '\0';
63     unsigned long long hash = 5381;
64     int c;
65
66     while ((c = *str++)) {
67         hash = ((hash << 5) + hash) + c;
68     }
69     return hashInt(hash, 8701123, seed) % mod;
70 }
71
72 // given a record and the field of interest, hashes it and returns a value
73 template<typename T> int hashRecord(std::string seed, T rec, unsigned int mod, std::string
    const& field)
74 {
75     unsigned int s = hashString(seed, 8701123, 0);
76     std::string test = rec.getCol(field);
77     return hashString(rec.getCol(field), mod, s);
78 }
79
80 // frees memory allocated for a hash index
81 void destroyHashIndex(linkedRecordPtr **hashIndex, unsigned int size);
82
83 #endif
84
```

```
1  /*
2  * DBMS Implementation
3  */
4  #include "stdafx.h"
5  #include "recordOps.h"
6  #include <stdlib.h>
7
8  // frees the memory allocated to a hash index
9
10 void destroyHashIndex(linkedRecordPtr **hashIndex, unsigned int size) {
11     for (unsigned int i = 0; i < size; i++) {
12         if (!hashIndex[i]) {
13             continue;
14         }
15         linkedRecordPtr *temp1 = hashIndex[i];
16         linkedRecordPtr *temp2 = temp1->next;
17         while (temp1) {
18             free(temp1);
19             temp1 = temp2;
20             if (temp2) {
21                 temp2 = temp2->next;
22             }
23         }
24     }
25     free(hashIndex);
26 }
27
```